#### Financial Data Science

# Lecture 8 Deep Reinforcement Learning

Video tutorial:

https://www.youtube.com/watch?v=Yd3H\_h-7C58

Liu Peng

liupeng@smu.edu.sg

# Introduction to Deep Reinforcement Learning (DRL)

- DRL combines deep neural networks (DNNs) with reinforcement learning (RL) principles to enable agents to learn optimal actions through interaction with the environment.
- Agent: Learns and makes decisions.
- Environment: Context or system where the agent operates.
- Policy: Strategy the agent follows to select actions.
- Reward Signal: Feedback to reinforce good actions.
- Goal: Agents aim to maximize cumulative rewards by adjusting their actions through trial-and-error exploration.

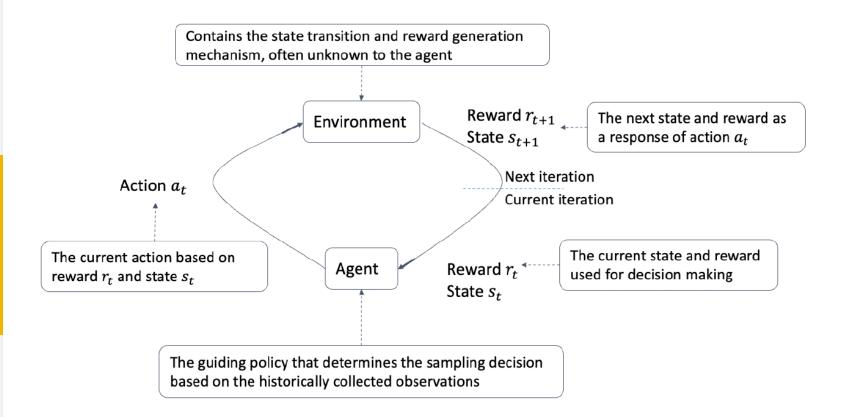


Figure 1: Iterative interaction between the agent and the environment.

### Interaction between Agent and Environment

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \cdots$$
$$G_t = R_{t+1} + \gamma G_{t+1}$$

$$V^{\pi}(s) = \mathbb{E}_{\pi}[G_t|S_t = s]$$

$$V^{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1}|S_t = s] + \gamma \mathbb{E}_{\pi}[G_{t+1}|S_t = s]$$

$$V^{\pi}(s) = \sum_{a} \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V^{\pi}(s')]$$

$$Q^{\pi}(s, a) = \mathbb{E}_{\pi}[G_t|S_t = s, A_t = a]$$

$$Q^{\pi}(s, a) = \sum_{s', r} p(s', r|s, a)[r + \gamma \sum_{a'} \pi(a'|s')Q^{\pi}(s', a')]$$

# DRL for Portfolio Optimization

#### **Objective:**

• Dynamically allocate assets to maximize returns, manage risk, or achieve specific financial goals through learned investment policies.

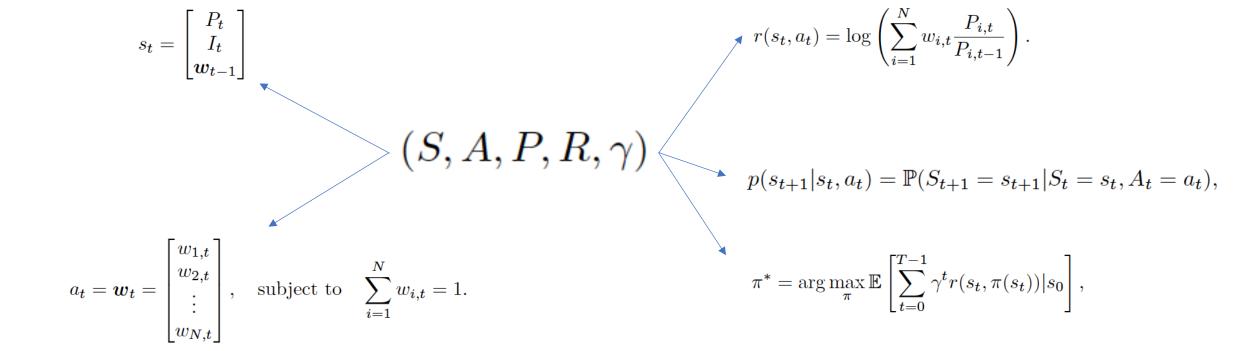
#### **DRL Formulation:**

- Agent: Portfolio manager making sequential investment decisions.
- Environment: Market conditions represented by financial indicators.
- Actions: Adjustments to asset weights within a portfolio.
- **Reward:** Risk-adjusted returns, e.g., Sharpe ratio or cumulative returns minus transaction costs.

#### **Key Advantages:**

- Adapts to changing market environments.
- Captures complex, nonlinear market dynamics.
- Integrates various constraints (e.g., risk, transaction costs, liquidity).

## DRL in Portfolio Optimization



## DRL for Hedging

**Goal:** Manage risk in derivative transactions (e.g., options).

**Traditional Approach:** Delta hedging (BSM model). Assumes frictionless, continuous-time markets.

#### **Real-World Problems:**

- **Discrete Time:** Hedging occurs periodically (e.g., daily), not continuously.
- Transaction Costs: Buying/selling the underlying asset incurs costs.
- Imperfect Replication: Leads to tracking errors and mis-hedging risk.
- Need: A strategy that balances minimizing hedging errors and reducing transaction costs, adapting to market dynamics.

## Applying RL to Derivative Hedging

Framework: Model the hedging problem as a Markov Decision Process.

State ( $s_t$ ): Can include:

- Time to maturity (t)
- Underlying asset price  $(S_t)$
- Current hedge position  $(N_t)$
- Option price  $(C_t)$ , Delta  $(\Delta_t)$  (optional, sometimes learned)
- Strike price (*K*)

Action ( $a_t$ ): The new hedge position ( $N_{t+1}$ ) or the change in position ( $\delta N_t$ ). Can be discrete or continuous.

**Reward** ( $r_t$ ): Designed to achieve the hedging objective. Common forms:

- Minimize P&L variance: Penalize the difference between the hedge portfolio's value change and the option's value change.
- Mean-Variance Optimization:  $r_t=\mathbb{E}[\delta\Pi_t]-\frac{\lambda}{2}Var[\delta\Pi_t]$  (Balances expected P&L change and its variance, factoring in transaction costs).
- Accounting P&L: Immediate reward based on step-by-step change in net portfolio value.

## Key RL Methods in Derivative Hedging

#### QLBS (Halperin, 2017):

- Applied Q-learning to option hedging in a Black-Scholes world.
- Minimized terminal variance.
- Used discrete states/actions, no transaction costs initially.

#### Deep Hedging (Buehler et al., 2019):

- Used neural networks as function approximators.
- Incorporated transaction costs and convex risk measures.

#### Automated Hedging (Ritter & Kolm, 2019):

- RL agent learns optimal strategy considering transaction costs.
- Used Q-learning/SARSA with continuous states, discrete actions.
- Showed RL comparable to delta hedging variance but lower cost.

#### Advanced Algorithms (Du et al., 2020):

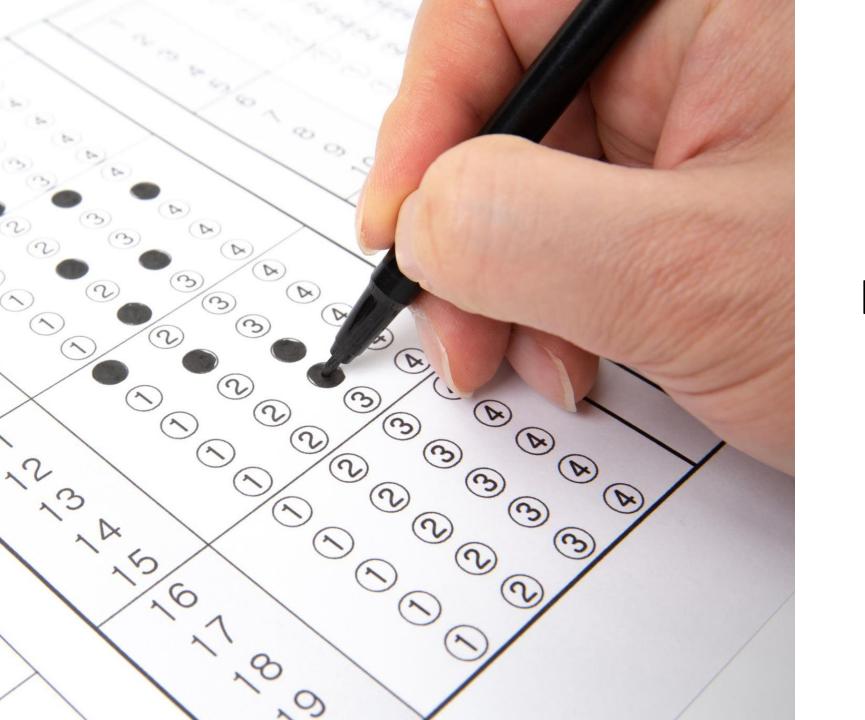
- Applied DQN, PPO to learn policy/action function directly.
- PPO showed strong performance (delta neutrality, training time/data).
- Learned option price/delta implicitly.

#### Continuous Actions (Cao et al., 2021):

- Used DDPG for continuous state and action spaces, reducing discretization error.
- Introduced decoupled reward (expectation & std dev of cost).
- Favored accounting P&L for faster learning.
- Outperformed traditional delta hedging with transaction costs & stochastic volatility.

#### Beyond Delta (Cao et al., 2023):

- Hedging Gamma & Vega using options as instruments.
- Used Distributional RL for more nuanced risk management (VaR, CVaR).
- Included Greeks in the state space.



In-class quiz