

QF632-2025-W1

Number of participants: 41



1. Which of the following is an example of unstructured data?

37 correct answers
out of 40 respondents

A CSV file containing tabular sales records

0%

0 votes

A JSON file with user demographic information

8%

3 votes

A collection of tweets without a predefined format

93%

37 votes

A spreadsheet of financial numbers

0%

0 votes




2.

What is the primary distinguishing feature of a dataframe compared to a generic 2D matrix?

34 correct answers
out of 37 respondents

Dataframes can be used only for storing Boolean values.



1 vote

Dataframes allow labeled rows and columns, enabling mixed data types in different columns.




34 votes

Dataframes store data only in numeric form without any text labels.



0 votes

Dataframes must be 3D, making them unsuitable for flat data storage.



2 votes



3. Which of the following best describes a scenario for using a while loop instead of a for loop?

34 correct answers
out of 38 respondents

You have a finite list of numbers to process in sequence.

0%

0 votes

You know exactly how many times you need to repeat the loop beforehand.

5%

2 votes

You want to iterate until a certain condition is no longer satisfied, but you don't know how many iterations it will take in advance.



89%

34 votes

You need to iterate over an array in a strictly defined order until the array ends.

5%

2 votes



What is the defining 4. characteristic of a recursive function?

**25 correct
answers**
out of 39
respondents

It requires an
iterative loop
inside its body.



13 votes



It calls itself
within its own
definition to
solve a problem.



25 votes

It can only be
used to solve
mathematical
problems like
factorials.



1 vote

It is
automatically
optimized by the
compiler to run in
parallel.



0 votes



5.

A researcher wants to run time-series analysis on repeated measurements for each subject across multiple time points. Converting the dataset into long format would be advantageous because:

26 correct answers

out of 31 respondents

It makes it easier to create pivot tables for summarizing data.



2 votes



Each time point appears in a separate row, allowing easy grouping by time and subject.



26 votes

It automatically encodes categorical variables as integers.



2 votes

It reduces memory usage compared to wide format.



1 vote



6. What is the primary purpose of performing one-hot encoding on a categorical variable?

33 correct answers
out of 34 respondents

To reduce the dimensionality of the dataset.

0%

0 votes

To represent each category with a separate binary feature (column), thus making it numeric.



97%

33 votes

To convert text data into continuous numeric ranges.

0%

0 votes

To compress continuous variables into a single categorical feature.

3%

1 vote



7.

Which statement best summarizes the difference between one-hot encoding and label encoding?

25 correct answers

out of 31 respondents

One-hot encoding is used for numeric data, while label encoding is used for text data only.




3 votes

One-hot encoding assigns each category a unique integer, while label encoding uses a binary column for each category.



3 votes

One-hot encoding creates multiple binary features for each category, whereas label encoding replaces each category with a single integer label.



25 votes

Both methods always reduce the overall number of



0 votes

columns needed
in the dataset.



Why is it important to split 8. the dataset into training, validation, and test sets?

0 correct answer
out of 0
respondent

To increase the
overall size of
the dataset for
training

0%

0 votes

To ensure that
performance
metrics are
estimated on
unseen data and
to reduce
overfitting



0%

0 votes

To reduce the
computational
cost of the
training process

0%

0 votes

To merge
different
datasets for a
larger training
set

0%

0 votes



9. What does hyperparameter tuning involve in the context of training a supervised model?

0 correct answer
out of 0
respondent

Adjusting the internal model weights during training by minimizing the loss function

0%

0 votes

Searching for the best configuration of settings (like learning rate, regularization strength, or network architecture) that control the training process



0%

0 votes

Evaluating the final model on the test dataset

0%

0 votes

Preprocessing the dataset for better performance

0%

0 votes



10. Which property is most essential for a loss function to be effectively minimized using gradient descent?

0 correct answer
out of 0
respondent

Boundedness
between 0 and 1

0%

0 votes



Differentiability

0%

0 votes

Invariance to
feature scaling

0%

0 votes

Non-convexity

0%

0 votes



11.

Which loss function is designed to combine the robustness of MAE with the smooth differentiability of MSE, making it less sensitive to outliers?

0 correct answer
out of 0
respondent

Log-Cosh Loss

0%

0 votes



Huber Loss

0%

0 votes

Cross-Entropy
Loss

0%

0 votes

Quantile Loss

0%

0 votes