

QF632-2025-W7

Number of participants: 24



1. **How many hidden layers does a three-layer ANN contain?**

16 correct answers
out of 17 respondents

1



1 vote



2



16 votes

3



0 votes



2. **During training, you observe that if you increase the batch size from 32 to 256, the training loss decreases more smoothly but convergence slows down. You decide to run experiments with batch sizes 16,, 32,, 64,, 128,, 256. This process is:**

14 correct answers
out of 18 respondents

Parameter tuning, because batch size affects the mini-batch gradients that update weights.



1 vote



Hyperparameter tuning, because batch size controls the training process but isn't a weight or bias.



14 votes

Neither, because batch size only affects hardware utilization, not learning.



0 votes

Both, since batch size influences gradient variance (parameter updates) and training speed (hyperparameter effect).



3 votes



3.

Which of the following statements is always true about a single-layer neural network (no hidden layers) with a linear activation function compared to ordinary least squares (OLS) linear regression?

11 correct answers
out of 16 respondents

It can model interaction terms (e.g., products of input features) if weights are tuned appropriately.



31%

5 votes

Its loss surface (mean-squared error) is non-convex in the weights.



0%

0 votes



It yields exactly the same fitted predictions as OLS linear regression when trained under MSE loss.



69%

11 votes

It can represent polynomial relationships between inputs and outputs without changing the architecture.



0%

0 votes



4.

Backpropagation means we take the input values and calculate the output of all layers until we reach the output layer

12 correct answers
out of 14 respondents

True



2 votes



False



12 votes



5.

What is the output value of node, with 3 links: H0: 1, w0: 0.1 | H1: 2, w1: 0.2 | H2: 3, w2: 0.3, & sigmoid activation?

13 correct answers
out of 16 respondents

 $0.1*1+0.2*2+0.3*3$ 

0 votes

 $\frac{-}{(0.1*1+0.2*2+0.3*3)}$ 

0 votes

 $\frac{1}{e^{(0.1*1+0.2*2+0.3*3)}}$ 

3 votes

 $\frac{1}{e^{[-(0.1*1+0.2*2+0.3*3)']}}$ 

13 votes



What happens when we increase 6. the number of hidden layers in a neural network?

16 correct answers
out of 17 respondents

The model
becomes more
prone to
underfitting.

0%

0 votes

The model
becomes more
prone to
overfitting.

94%

16 votes

The model's
training time
decreases.

6%

1 vote

The model
becomes less
capable of learning
complex patterns.

0%

0 votes



7. What is the gradient descent's objective?

12 correct answers
out of 13 respondents

Find a maximum
for the loss
function

0%

0 votes

Find a solution for
'loss function = 0'

8%

1 vote



Find a minimum for
the loss function

92%

12 votes



8. Which application is not suitable for RNN?

3 correct answers
out of 15 respondents

Score audio recording of IELTS takers



4 votes

Predicting default with 1000 customers' credit default data



3 votes

Predict stock movement using stock price of Apple over 10 years



5 votes

Translate a paragraph from English to Malay



3 votes



9. Which of the following can be a challenge when using SGD?

12 correct answers
out of 12 respondents

Choosing an appropriate learning rate.

0%

0 votes

Dealing with noisy gradients.

0%

0 votes



Both a and b.

100%

12 votes

Neither a nor b.

0%

0 votes



10. What is the main advantage of SGD over Batch Gradient Descent?

8 correct answers
out of 12 respondents



SGD is computationally less expensive.



8 votes

SGD always finds the global minimum.



4 votes

SGD uses more memory.



0 votes

There is no advantage of SGD over Batch Gradient Descent.



0 votes