

# Machine Learning and Financial Applications

## Lecture 4

### Deep Reinforcement Learning in Portfolio Optimization

---

Liu Peng

liupeng@smu.edu.sg

Video tutorial

[https://youtu.be/Yd3H\\_h-7C58?si=\\_gHkcvTs-TePQdUB](https://youtu.be/Yd3H_h-7C58?si=_gHkcvTs-TePQdUB)

# Learning outcomes



Get to know RL  
fundamentals



Able to create RL  
environment for portfolio  
optimization



Able to train and evaluate  
RL agent that performs  
portfolio optimization

# RL Learning Framework

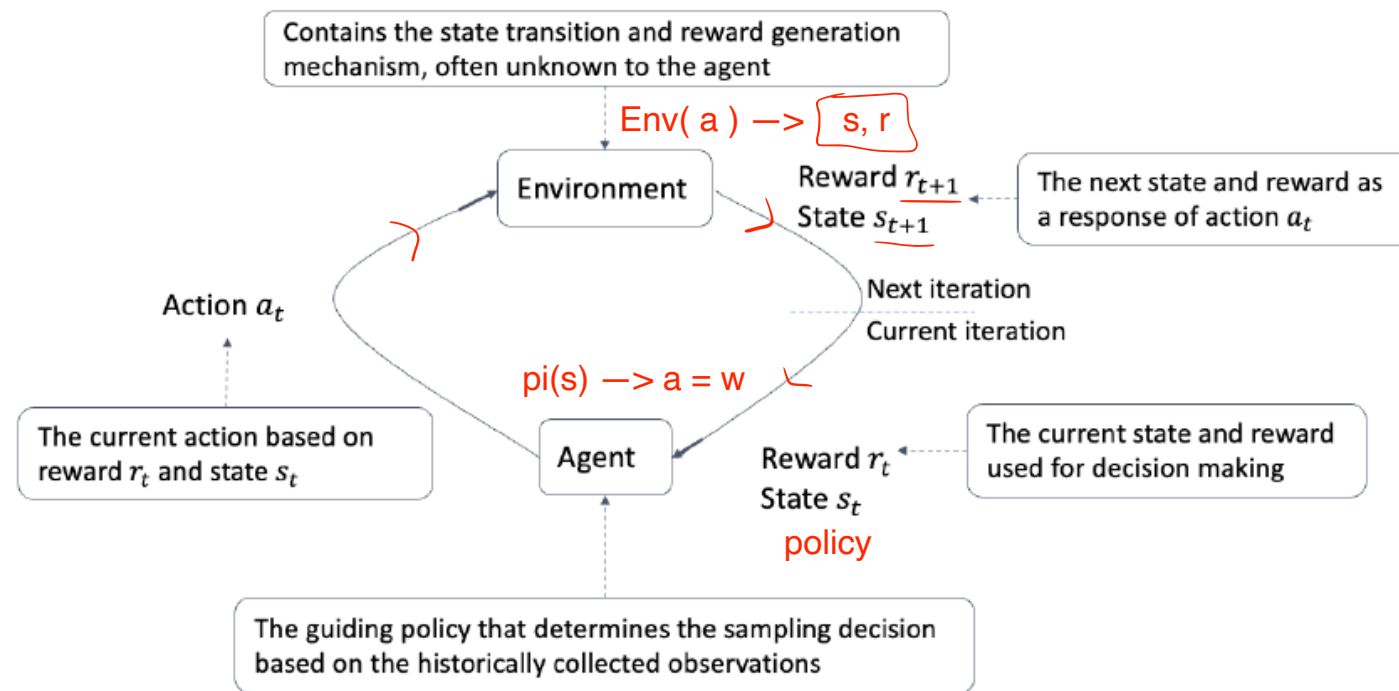


Figure 1.1: Iterative interaction between the agent and the environment.

# Using Reinforcement Learning for Portfolio Optimization

## Agent–Environment Setup

- State: market observations (prices, indicators, macro factors)
- Action: portfolio weight vector allocation

## Reward Design

- Risk-adjusted return

## Key Architectures

- Value-based: Deep Q-Networks (DQN)
- Policy-based: PPO, A2C/A3C
- Actor-Critic: DDPG, SAC

## Practical Considerations

- Transaction costs & market impact
- Position-sizing constraints (e.g. leverage, turnover limits)

## Stability & Efficiency

- Experience replay, target networks
- Off-policy vs. on-policy trade-offs

## Risk Management Integration

- Constraint handling via penalty in reward or Lagrangian methods
- Incorporating CVaR/drawdown objectives


# Review of Modern Portfolio Theory

LR is good at PO, bad at forecasting

$$\begin{array}{ccc}
 \underbrace{\text{minimize}}_{\mathbf{w}} & \underline{\mathbf{w}^T \Sigma \mathbf{w}} \text{ ptf risk} & \longleftrightarrow \text{maximize}_{\mathbf{w}} \quad \mathbf{w}^T \hat{\boldsymbol{\mu}} \uparrow \\
 \text{subject to} & \begin{cases} \mathbf{w}^T \boldsymbol{\mu} = \underline{\mu_0}, \text{ ptf target return} \\ \mathbf{w}^T \mathbf{1} = 1, (\#eq : \text{meanvar}) \end{cases} & \text{subject to} \quad \begin{cases} \mathbf{w}^T \Sigma \mathbf{w} \leq \underline{\sigma_p^2}, \\ \mathbf{w}^T \mathbf{1} = 1, \end{cases}
 \end{array}$$

$$\begin{array}{ccc}
 \downarrow & \mathcal{L}(\mathbf{w}, \lambda_1, \lambda_2) = \mathbf{w}^T \Sigma \mathbf{w} - \lambda_1 (\mathbf{w}^T \boldsymbol{\mu} - \mu_0) - \lambda_2 (\mathbf{w}^T \mathbf{1} - 1). \star & \uparrow \\
 \rightarrow & \underline{\frac{\partial \mathcal{L}}{\partial \mathbf{w}}} = 2\Sigma \mathbf{w} - \lambda_1 \boldsymbol{\mu} - \lambda_2 \mathbf{1} = \underline{0}. & \leftarrow
 \end{array}$$

$$\underline{\mathbf{w}^*} = \left( \frac{C\mu_0 - B}{AC - B^2} \right) \Sigma^{-1} \boldsymbol{\mu} + \left( \frac{A - B\mu_0}{AC - B^2} \right) \Sigma^{-1} \mathbf{1}$$



# In-class quiz

- Q1-4

# Group Discussion

- Propose a few user cases for RL based on your past project/work experiences

# Why Deep RL for PO?

---

## **Model-Free End-to-End Learning**

- Learns trading policies directly from raw market data, bypassing explicit return/risk prediction steps.

## **Continuous Adaptation**

- Updates strategies online as new data arrive, reducing the need for periodic retraining on historical samples.

## **Capturing Non-Linear Dependencies**

- Leverages deep networks to model complex asset relationships for more informed allocation decisions.

## **Scalable to High Dimensions**

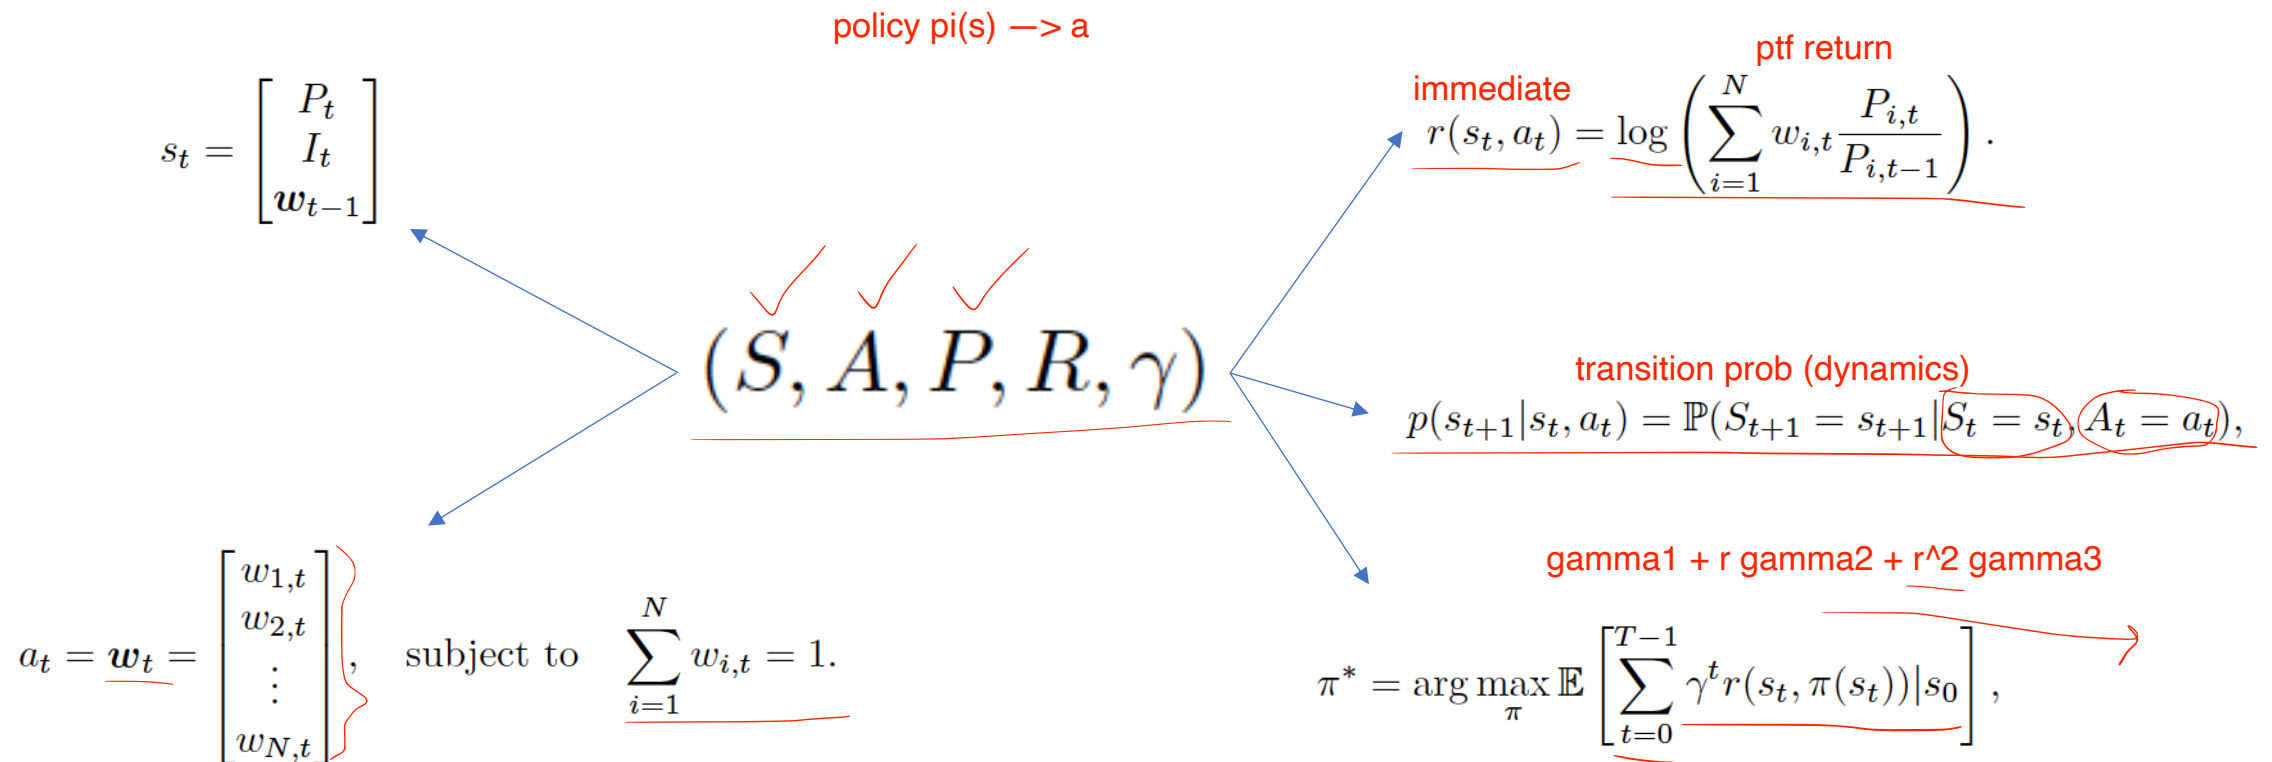
- Handles large universes of assets and intricate market structures, making it suitable for real-world portfolio tasks.

## **Built-in tradeoff between exploration and exploitation**



# Markov Decision Process

- A mathematical framework that provides a formal description of an environment for RL.



# RL Environment Setup for PO

- **State Space** ( $\mathcal{S}$ ): Encodes market information at time  $t$ :

$$s_t = \begin{bmatrix} P_t \\ I_t \\ \mathbf{w}_{t-1} \end{bmatrix}, \quad \text{where } P_t : \text{price data, } I_t : \text{market indicators, } \mathbf{w}_{t-1} : \text{previous weights.}$$

- **Action Space** ( $\mathcal{A}$ ): Portfolio allocation weights at time  $t$ :

$$a_t = \mathbf{w}_t = \begin{bmatrix} w_{1,t} \\ w_{2,t} \\ \vdots \\ w_{N,t} \end{bmatrix}, \quad \text{subject to } \sum_{i=1}^N w_{i,t} = 1.$$

- **Reward Function** ( $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ ): Immediate payoff (log portfolio return):


$$r(s_t, a_t) = \log \left( \sum_{i=1}^N w_{i,t} \frac{P_{i,t}}{P_{i,t-1}} \right).$$

- **Transition Dynamics** ( $p : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ ): Probability of transitioning to next state, implicitly learned:

$$p(s_{t+1} | s_t, a_t) = \mathbb{P}(S_{t+1} = s_{t+1} \mid S_t = s_t, A_t = a_t).$$

- **Learning Objective**: Find optimal policy  $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$  maximizing cumulative discounted rewards:

$$\pi^* = \arg \max_{\pi} \mathbb{E} \left[ \sum_{t=0}^{T-1} \gamma^t r(s_t, \pi(s_t)) \right], \quad \gamma \in [0, 1].$$



# In-class quiz

- Q5-8

# Value Functions

- **State-Value Function:** Expected cumulative future reward starting from state  $s$  under policy  $\pi$ :

$$V^\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s].$$

- **Action-Value Function:** Expected cumulative future reward starting from state  $s$ , taking action  $a$ , then following policy  $\pi$ :

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a].$$

- **Long-Term Return ( $G_t$ ):** Sum of discounted future rewards, expressed recursively:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = R_{t+1} + \gamma G_{t+1},$$

where  $\gamma \in [0, 1]$  is a discount factor, balancing immediate vs. future rewards.

- **Interpretation:** Value functions quantify the expected performance (long-term returns) of a given policy, guiding the selection and improvement of decision-making strategies.

# Bellman Equations for State and Action-Value Functions

- **Bellman Equation (Intuition):** Defines a recursive relationship between the value of a state (or state-action pair) and its successor states (or state-action pairs), based on the principle of optimality.
- **State-Value Bellman Equation:** Expresses value of a state as the immediate reward plus discounted future rewards:

$$V^\pi(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V^\pi(s')] .$$

- **Action-Value Bellman Equation:** Expresses value of state-action pair as the immediate reward plus discounted future rewards:

$$Q^\pi(s,a) = \sum_{s',r} p(s',r|s,a) \left[ r + \gamma \sum_{a'} \pi(a'|s') Q^\pi(s',a') \right] .$$

- **Significance in RL:** These recursive equations form the foundation for iterative algorithms (value iteration, policy iteration, Q-learning) and are fundamental in deep RL for training value-function approximators.

# Optimal Value Functions

The optimal state-value function  $v^*(s)$  is the maximum value function over all policies:

$$v^*(s) = \max_{\pi} v_{\pi}(s). \quad (11)$$


The optimal action-value function  $q^*(s, a)$  is the maximum action-value function over all policies:

$$q^*(s, a) = \max_{\pi} q_{\pi}(s, a). \quad (12)$$

The Bellman optimality equations are:

$$v^*(s) = \max_{a \in A} \mathbb{E}[\overset{\text{immediate reward}}{R_{t+1}} + \gamma \overset{\text{future}}{v^*(S_{t+1})} | \underset{\text{signal}}{S_t = s, A_t = a}], \quad (13)$$

$$q^*(s, a) = \mathbb{E}[R_{t+1} + \gamma \max_{a'} q^*(S_{t+1}, a') | S_t = s, A_t = a]. \quad (14)$$



# In-class quiz

- Q9-12

# Q-Learning

---

**Algorithm Type:** Off-policy, model-free, value-based RL for learning  $Q^*(s, a)$  without knowing dynamics.

**Objective:** Estimate

$$Q^*(s, a) = \max_{\pi} \mathbb{E}[G_t \mid S_t = s, A_t = a], \quad \pi^*(s) = \arg \max_a Q^*(s, a).$$

**Core Update Rule:**

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ \overbrace{r + \gamma \max_{a'} Q(s', a')}^{\text{long-term}} - Q(s, a) \right],$$

where  $\alpha$  is the learning rate and  $\gamma$  the discount factor.

**Exploration–Exploitation:** Use  $\varepsilon$ -greedy:

- With probability  $\varepsilon$ , choose a random action.
- Otherwise, choose  $a = \arg \max_a Q(s, a)$ .

**Convergence Guarantees:** With decaying  $\{\alpha_t\}$  (s.t.  $\sum_t \alpha_t = \infty$ ,  $\sum_t \alpha_t^2 < \infty$ ) and sufficient exploration,  $Q(s, a) \rightarrow Q^*(s, a)$ .

**Practical Considerations:**

- *Tabular vs. Function Approximation:* Tabular for small discrete spaces; use neural nets when large/continuous.
- *Stability Tricks:* Experience replay and target networks (in Deep Q-Learning) to decorrelate updates.



# Deep Q Network (DQN)

## Core Idea:

Uses deep neural networks to approximate the Q-value function:

$$Q(s, a; \theta) \approx Q^*(s, a)$$

## Algorithmic Components:

- **Experience Replay:**  
Stores past experiences  $(s, a, r, s')$  to stabilize training.
- **Target Network:**  
Periodically updated copy of Q-network to reduce correlations and improve stability.
- **Epsilon-Greedy Policy:**  
Balances exploration (random action) and exploitation (best-known action).

## Loss Function (Bellman Error):

Minimizes difference between current and target Q-values:

$$L(\theta) = \mathbb{E} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right]$$

## Applications in Portfolio Optimization:

Learns optimal trading policies directly from market states without explicit forecasting of returns or risks.

# Implementing DQN




Exhibit 1: Daily cumulative returns of GOOG, MSFT, AAPL, and TSLA.




Exhibit 2: Cumulative wealth over time during the training period for all portfolio strategies. DQN and Double DQN performed best, followed by Dueling DQN and equal weighted portfolio. GMVP and random policy performed worst.

# Deep Deterministic Policy Gradient (optional)

## Actor-Critic Architecture:

- **Actor:** Represents the deterministic policy  $\pi_\theta(s)$  that maps states  $s$  directly to continuous actions  $a$ .
- **Critic:** Estimates the action-value function  $Q_\phi(s, a)$ , evaluating the expected return of action  $a$  in state  $s$ .

## Deterministic Policy Gradient Theorem:

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{s \sim \mathcal{D}} \left[ \nabla_\theta \pi_\theta(s) \nabla_a Q^\pi(s, a) \Big|_{a=\pi_\theta(s)} \right]$$

See accompanying notes for derivation

- Guides the Actor to adjust policy parameters  $\theta$  in directions that maximize the expected return.

## Continuous Action Spaces:

- Enables precise control over actions without discretization, addressing scalability and precision issues inherent in traditional Q-learning.

# Enhancements for Efficient Learning in DDPG (optional)

## Experience Replay Buffer ( $\mathcal{D}$ ):

$$\mathcal{D} = \{(s_t, a_t, r_t, s_{t+1})\}$$

- Stores past experiences to break temporal correlations and improve sample efficiency through random minibatch sampling.

## Off-Policy Learning:

- Enables learning from historical data irrespective of the current policy, facilitating effective exploration and exploitation.

## Actor-Critic Collaboration:

- **Critic** provides accurate  $Q$ -value estimates to guide the **Actor**.
- **Actor** continuously improves the policy to select actions that maximize the Critic's estimated returns, ensuring coherent and stable learning dynamics.

# Actor-Critic Architecture (optional)

**Critic Update via Bellman Equation:**

$$L(\phi) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}} \left[ \left( Q_{\phi}(s, a) - (r + \gamma Q'_{\phi'}(s', \pi'_{\theta'}(s'))) \right)^2 \right]$$

- Minimizes the Mean Squared Bellman Error (MSBE) to accurately estimate  $Q$ -values.

**Actor Update Using Policy Gradient:**

$$\theta \leftarrow \theta + \alpha_{\theta} \nabla_{\theta} J(\pi_{\theta})$$

- Applies gradient ascent to optimize the policy based on Critic feedback.

**Target Networks for Stability:**

$$\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$$

$$\theta' \leftarrow \tau \theta + (1 - \tau) \theta'$$

- Soft updates with factor  $\tau \ll 1$  ensure gradual parameter changes, mitigating oscillations and divergence during training.

# DDPG Algorithm (optional)

## Unique advantages

1. **Initialize** Actor, Critic, Target Networks, and Replay Buffer.

2. **For each episode:**

- Initialize state  $s_0$ .
- **For each timestep:**
  - Select action  $a_t = \pi_{\theta}(s_t) + \mathcal{N}_t$ .
  - Execute  $a_t$ , observe  $r_t$  and  $s_{t+1}$ .
  - Store  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$ .
  - Sample minibatch from  $\mathcal{D}$ .
  - Update Critic and Actor networks.
  - Soft update Target Networks.
- **Handles Continuous Action Spaces:** Directly outputs precise actions without discretization.
- **Sample Efficiency:** Utilizes Experience Replay to make better use of collected data.
- **Stable Learning:** Target Networks and soft updates prevent oscillations and divergence.
- **Off-Policy Learning:** Learns from past experiences regardless of the current policy, enhancing exploration.

```
mirror_mod = modifier_ob.  
#set mirror object to mirror  
mirror_mod.mirror_object  
operation == "MIRROR_X":  
mirror_mod.use_x = True  
mirror_mod.use_y = False  
mirror_mod.use_z = False  
operation == "MIRROR_Y":  
mirror_mod.use_x = False  
mirror_mod.use_y = True  
mirror_mod.use_z = False  
operation == "MIRROR_Z":  
mirror_mod.use_x = False  
mirror_mod.use_y = False  
mirror_mod.use_z = True
```

```
#selection at the end -add  
mirror_ob.select= 1  
modifier_ob.select=1  
context.scene.objects.active  
("Selected" + str(modifier_ob))  
mirror_ob.select = 0  
= bpy.context.selected_object  
data.objects[one.name].select  
print("please select exactly
```

--- OPERATOR CLASSES ---

```
types.Operator):  
on X mirror to the selected  
object.mirror_mirror_x"  
mirror X"
```

```
context):  
context.active_object is not
```

# Coding session

---

# Group Homework

- Open exercise
- Experiment with other types of RL algorithms, such as policy networks and actor-critic networks
- Consider other constraints and reward designs

Due one day before class starts next week





# Homework



Watch/review video tutorials  
and class recording



Post learning reflections and  
questions if any



Complete group homework